

Predicting Loan Default for a Bank

Name: Thomas Sugg G Number: G*****

```
# Add all library you will need here
```

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyver t.
```

```
## v ggplot2 3.1.0      v purrr  0.3.0
## v tibble  2.0.1      v dplyr  0.7.8
## v tidyr   0.8.2      v stringr 1.3.1
## v readr   1.3.1      v forcats 0.3.0
```

```
## -- Conflicts ----- tidyver
```

```
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
```

```
library(randomForest)
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##      combine
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
##      margin
```

```
library(factoextra)
```

```
## Welcome! Related Books: `Practical Guide To Cluster Analysis in R` at https://goo.gl/13EFCZ
```

```
library(corrplot)
```

```
## corrplot 0.84 loaded
```

```
library(rpart)
```

```
library(rpart.plot)
```

```
library(kknn)
```

```
library(MASS)
```

```
##
```

```
## Attaching package: 'MASS'
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##      select
```

```
library(ISLR)
```

```
# This will read in the data frame
```

```
loan_data <- readRDS(file = "/cloud/project/Final Project/loan_data.rds")
```

```

# Create training and test data
set.seed(314)
train_index <- sample(1:nrow(loan_data), floor(0.7*nrow(loan_data)))

# training
loan_training <- loan_data[train_index, ]

# test
loan_test <- loan_data[-train_index, ]

# Function for analyzing confusion matrices
cf_matrix <- function(actual_vec, pred_prob_vec, positive_val,
                      cut_prob = 0.5, search_cut = FALSE) {

  if (search_cut == FALSE) {
    actual <- actual_vec == positive_val; pred <- pred_prob_vec >= cut_prob
    P <- sum(actual); N <- length(actual) - P; TP <- sum(actual & pred)
    FN <- P - TP; TN <- sum(!(actual) & !(pred)); FP <- N - TN

    if (TP != 0) { Precision <- TP/(TP + FP); Recall <- TP/(TP + FN)
      F1 <- 2*((Precision*Recall)/(Precision + Recall))}

    if(TP == 0) { Precision = 0; Recall = 0; F1 = 0 }

    model_results <- list(confusion_matrix =
      data.frame(metric = c("Correct", "Misclassified", "True Positive",
        "True Negative", "False Negative", "False Positive"),
        observations = c(TN + TP, FN + FP, TP, TN, FN, FP),
        rate = c((TN + TP)/(N + P), (FN + FP)/(N + P), TP/P, TN/N, FN/P, FP/N),
        pct_total_obs = c((TN + TP), (FN + FP), TP, TN, FN, FP)*(1/(N + P)),
        stringsAsFactors = FALSE),
      F1_summary =
        data.frame(metric = c("Precision", "Recall", "F1 Score"),
          value = c(Precision, Recall, F1),
          stringsAsFactors = FALSE))
    return(model_results) }

  if (search_cut == TRUE) {
    optimal_cut = data.frame(cut_prob = seq(0,1, by = 0.05),
      correct_rate = NA, F1_score = NA,
      false_pos_rate = NA, false_neg_rate = NA)

    for (row in (1:nrow(optimal_cut))) {
      actual <- actual_vec == positive_val
      pred <- pred_prob_vec >= optimal_cut$cut_prob[row]
      P <- sum(actual); N <- length(actual) - P
      TP <- sum(actual & pred); FN <- P - TP
      TN <- sum(!(actual) & !(pred)); FP <- N - TN

      if (TP != 0) { Precision <- TP/(TP + FP); Recall <- TP/(TP + FN)
        F1 <- 2*((Precision*Recall)/(Precision + Recall))}

      if(TP == 0) { Precision = 0; Recall = 0; F1 = 0 }
    }
  }
}

```

```

    optimal_cut[row, 2:5] <- c((TN + TP)/(N + P), F1, FP/N, FN/P)
  }
  return(optimal_cut)
}

```

Loan Data

The `loan_data` data frame contains information on 3-year loans that were originated in 2013 by a local bank for customers residing in the United States. The company is looking to see if it can determine the factors that lead to loan default and whether it can predict if a customer will eventually default on their loan at time of loan origination. The goal is to become better at identifying customers at risk of defaulting on their loans to minimize the bank's financial losses.

The dataset contains a mixture of applicant demographics (gender, age, residence, etc.), financial information (income, debt ratios, FICO scores, etc.), and applicant behavior (number of open accounts, historical engagement with the bank's products, number of missed payments, etc...)

Specifically, the broad questions that the bank is trying to answer include:

1. What are the factors that contribute to customers defaulting on their loans?
2. Is it possible to predict whether a customer will default on their loan? If so, how accurate are the predictions?
3. How many costly errors does the predictive model produce (customers classified as not defaulting, but eventually do)?

Exploratory Data Analysis Section

1. Does the loan amount and the loan proportion of income determine loan default?

Findings: No, loan amount and loan proportion of income do not appear to have an effect on loan default. For default "Yes" and "No", both scatter plots look the same.

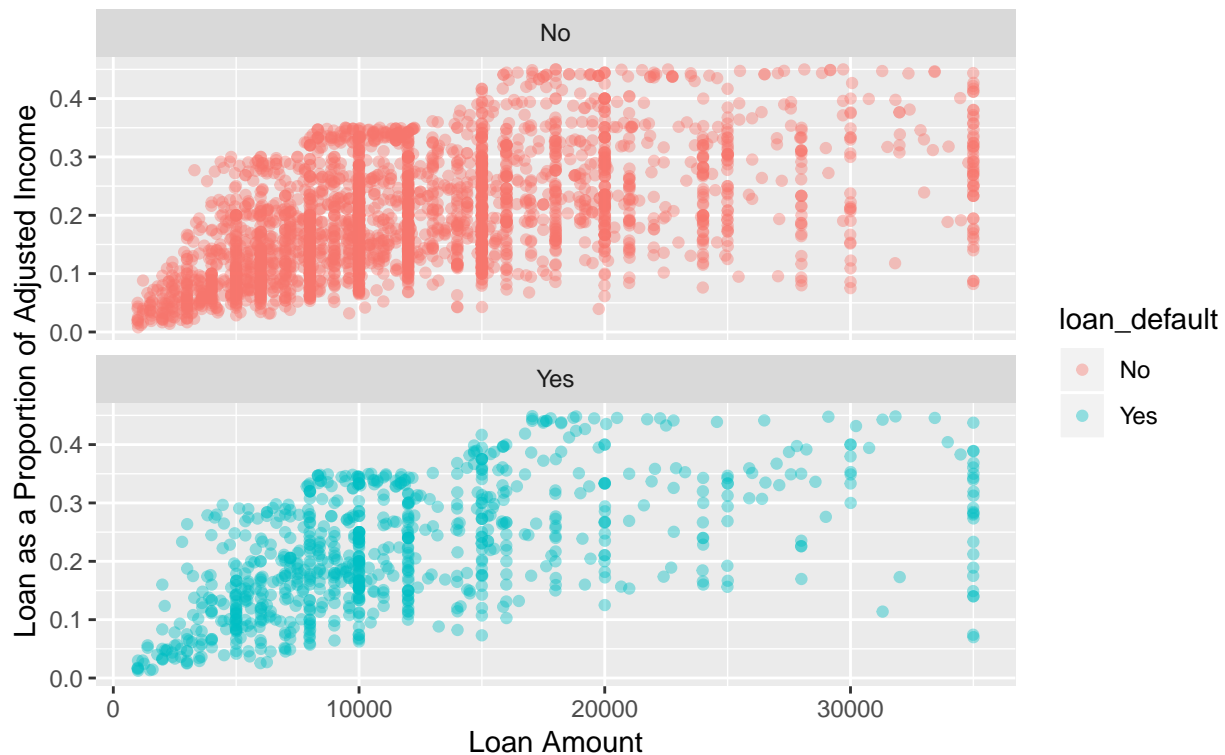
```

loan_amnt_graph <- ggplot(data = loan_data, mapping = aes(x = loan_amnt, y = pct_loan_income, color = loan_default)) +
  geom_jitter(alpha = 0.4) +
  facet_wrap(~loan_default, nrow = 2) +
  labs(title = "Loan Default Rates by Loan Amount
    and Loan Proportion of Income",
    x = "Loan Amount",
    y = "Loan as a Proportion of Adjusted Income")

loan_amnt_graph

```

Loan Default Rates by Loan Amount and Loan Proportion of Income



2. Do loan default rates differ by location of residence and type of residence?

Findings: Yes, customers who rent and own in the Northeast and Midwest have higher default rates than other customers. Default rates in these areas are more than double the rates in the West, Mid-Atlantic, South, and Southwest.

```
default_by_residence <- loan_data %>% group_by(us_region_residence, residence_property) %>% summarise(
  customers_who_defaulted = sum(loan_default == "Yes"),
  default_rate = customers_who_defaulted / total_customers
)
arrange(default_by_residence, desc(default_rate))
```

```
## # A tibble: 12 x 5
## # Groups:   us_region_residence [6]
##   us_region_residence residence_property total_customers customers_who_defaulted default_rate
##   <fct>               <fct>                <int>             <int>          <dbl>
## 1 Northeast          Rent                   327               155          0.474
## 2 Midwest             Rent                   189                81          0.428
## 3 Midwest             Own                    283               100          0.353
## 4 Northeast          Own                    385               133          0.345
## 5 West                Rent                   396                71          0.179
## 6 Mid-Atlantic        Rent                   439                75          0.171
## 7 South               Rent                   126                20          0.159
## 8 West                Own                    494                77          0.156
## 9 Southwest           Rent                   141                21          0.150
## 10 Mid-Atlantic       Own                    603                79          0.131
## 11 Southwest          Own                    198                23          0.116
## 12 South              Own                    164                12          0.073
## # ... with 1 more variable: default_rate <dbl>
```

3. Do loan default rates differ by education?

Findings: Yes, customers with the two lowest levels of education have the highest default rates. Interestingly, customers with the highest level of education have the third highest default rate.

```
default_by_education <- loan_data %>% group_by(highest_ed_level) %>% summarise(total_customers = n(),
                                         customers_who_defaulted = sum(loan_status == "D", na.rm = TRUE),
                                         default_rate = customers_who_defaulted / total_customers)

arrange(default_by_education, desc(default_rate))
```

```
## # A tibble: 5 x 4
##   highest_ed_level total_customers customers_who_defaulted default_rate
##   <fct>              <int>              <int>              <dbl>
## 1 High School         328                202              0.616
## 2 < High School       298                130              0.436
## 3 PhD or Doctorate   408                 94              0.230
## 4 Masters            877                138              0.157
## 5 Bachelors         1834                283              0.154
```

4. Do loan default rates differ by income bracket? Income levels determined by the Pew Research Center in 2017.

(lower = less than \$39,500). (middle = between \$39,500 and \$118,000). (upper = more than \$118,000).

Findings: Yes, default rates decrease as income increases.

```
income_levels <- cut(x = loan_data$adjusted_annual_inc,
                    breaks = c(-Inf, 39500, 118000, Inf),
                    labels = c("lower", "middle", "upper"),
                    right = TRUE)

income_levels_data <- cbind(loan_data, income_levels)

income_levels_data %>% group_by(income_levels) %>% summarise(total_customers = n(),
                                                             customers_who_defaulted = sum(loan_status == "D", na.rm = TRUE),
                                                             default_rate = customers_who_defaulted / total_customers)
```

```
## # A tibble: 3 x 4
##   income_levels total_customers customers_who_defaulted default_rate
##   <fct>              <int>              <int>              <dbl>
## 1 lower            1540                435              0.282
## 2 middle           1973                370              0.188
## 3 upper            232                 42              0.181
```

5. Is there an interaction between customer credit history and default rate? FICO score levels replicate those of Experian.

(very poor = between 300 and 579). (fair = between 580 and 669). (good = between 670 and 739). (very good = between 740 and 799). (exceptional = between 800 and 850).

Findings: Yes, lower FICO scores appear to have a higher default rate. The same can be said about credit inquiries.

```
fico_levels <- cut(x = loan_data$fico_score,
                  breaks = c(-Inf, 580, 670, 740, 800, 850),
                  labels = c("Very Poor", "Fair", "Good", "Very Good", "Exceptional"),
                  right = TRUE)

loan_fico_data <- cbind(loan_data, fico_levels)
```

```

default_by_fico <- loan_fico_data %>% group_by(fico_levels) %>% summarise(total_customers = n(),
                                                                           number_of_credit_inquiries = number_of_credit_inquiries,
                                                                           customers_who_defaulted = customers_who_defaulted,
                                                                           default_rate = customers_who_defaulted / total_customers)


default_by_fico

## # A tibble: 5 x 5
##   fico_levels total_customers number_of_credi~ customers_who_d~
##   <fct>          <int>          <int>          <int>
## 1 Very Poor      1339            1176            606
## 2 Fair           1571            1175            192
## 3 Good           653             497             42
## 4 Very Good      161             126              6
## 5 Exceptional     21              16              1
## # ... with 1 more variable: default_rate <dbl>

fico_default_graph <- ggplot(data = default_by_fico, mapping = aes(x = fico_levels, y = default_rate, fill = fico_levels)) +
  geom_bar(stat = "identity") +
  labs(title = "Loan Default Rates by FICO score",
       x = "FICO Score",
       y = "Default Rate")

fico_default_graph

```



6. Is there a relationship between gender and age that may predict default rates?

Findings: Yes, men consistently have higher default rates than women. As both genders become older, their default rates tend to decrease; however, default rates increase at the last two age categories.

```

default_by_gender <- loan_data %>% group_by(gender, age_category) %>% summarise(total_customers = n(),
                                                                 customers_who_defaulted = sum(loan_de
                                                                 default_rate = customers_who_defaulted

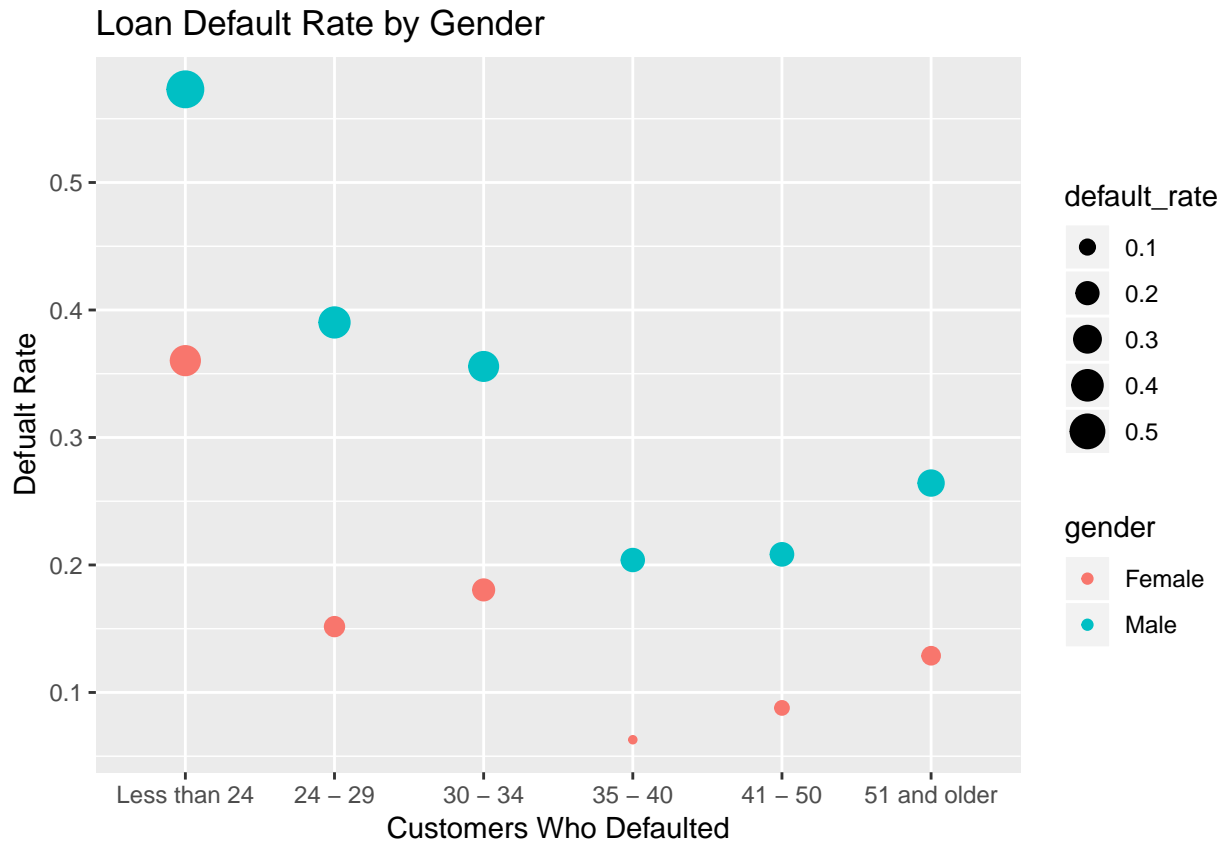
arrange(default_by_gender, desc(default_rate))

## # A tibble: 12 x 5
## # Groups:   gender [2]
##   gender age_category total_customers customers_who_defaulted default_rate
##   <fct>   <fct>           <int>           <int>           <dbl>
## 1 Male   Less than 24           260             149             0.573
## 2 Male   24 - 29                 287             112             0.390
## 3 Female Less than 24           297             107             0.360
## 4 Male   30 - 34             253              90             0.356
## 5 Male   51 and older        193              51             0.264
## 6 Male   41 - 50             264              55             0.208
## 7 Male   35 - 40             309              63             0.204
## 8 Female 30 - 34             266              48             0.180
## 9 Female 24 - 29             455              69             0.152
## 10 Female 51 and older        295              38             0.129
## 11 Female 41 - 50             421              37             0.0879
## 12 Female 35 - 40             445              28             0.0629

default_gender_plot <- ggplot(default_by_gender, mapping = aes(x = age_category, y = default_rate, color =
  geom_point(mapping = aes(size = default_rate))) +
  labs(title = "Loan Default Rate by Gender",
       x = "Customers Who Defaulted",
       y = "Default Rate")

default_gender_plot

```



7. Does the number of accounts 120 days overdue and public bankruptcies influence default rates?

Findings: The proportion of customers who were 120 days past due and the proportion of customers who had publicly filled bankruptcy were higher for customers who defaulted. However, the difference is very small and may prove unimportant in variable selection.

```
default_by_bankruptcies <- loan_data %>% group_by(loan_default) %>% summarise(number_of_customers = n(),
  num_accounts_open = sum(num_accounts_open),
  num_bankruptcies = sum(num_bankruptcies),
  prop_acct_overdue = num_accounts_open / number_of_customers,
  prop_bankruptcies = num_bankruptcies / number_of_customers)

default_by_bankruptcies
```

```
## # A tibble: 2 x 6
##   loan_default number_of_custome~ num_accounts_op~ num_bankruptcies
##   <fct>          <int>          <int>          <int>
## 1 No             2898             915             247
## 2 Yes             847             279             89
## # ... with 2 more variables: prop_acct_overdue <dbl>,
## #   prop_bankruptcies <dbl>
```

8. Does a customer's debt to income ratio influence default rates?


Findings: Customers who defaulted have a higher median debt to income ratio. However, it is only a small difference compared to those who did not default.

```
default_by_dti <- loan_data %>% group_by(loan_default)

default_by_dti
```



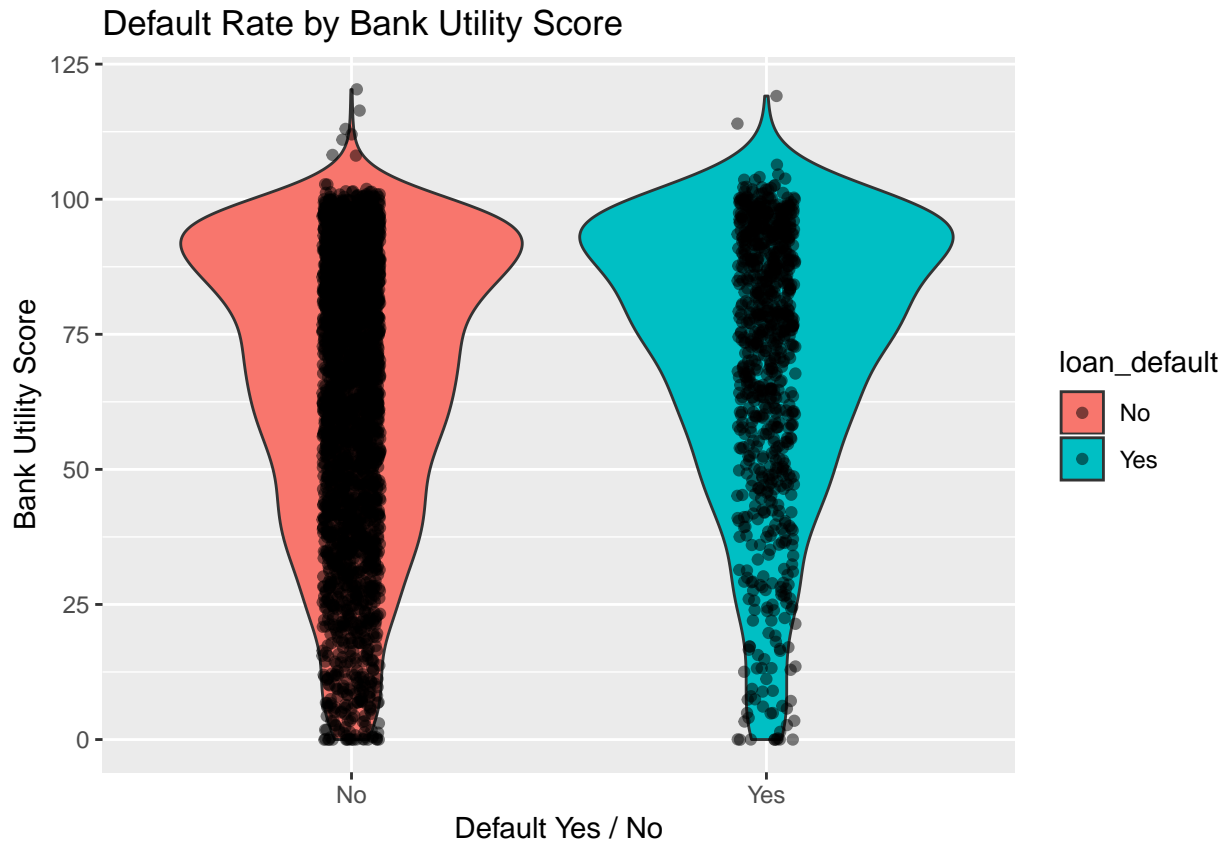
```
## # A tibble: 3,745 x 16
## # Groups:   loan_default [2]
##   loan_default residence_prope~ gender age_category highest_ed_level
##   <fct>         <fct>         <fct> <fct>         <fct>
## 1 No          Rent          Female 41 - 50    Bachelors
## 2 No          Own           Female 30 - 34    Bachelors
## 3 No          Own           Female 35 - 40    Bachelors
## 4 Yes         Rent          Female Less than 24 Bachelors
## 5 No          Own           Female 35 - 40    Masters
## 6 No          Own           Male   41 - 50    Bachelors
## 7 No          Own           Female 24 - 29    Bachelors
## 8 No          Own           Female 24 - 29    PhD or Doctorate
## 9 No          Own           Female Less than 24 Bachelors
## 10 No         Rent          Female 35 - 40    PhD or Doctorate
## # ... with 3,735 more rows, and 11 more variables:
## #   us_region_residence <fct>, loan_amnt <int>, adjusted_annual_inc <dbl>,
## #   pct_loan_income <dbl>, fico_score <dbl>, dti <dbl>,
## #   inq_last_6mths <int>, open_acc <int>, bc_util <dbl>,
## #   num_accts_ever_120_pd <int>, pub_rec_bankruptcies <int>
default_dti_plot <- ggplot(data = default_by_dti, mapping = aes(x = reorder(loan_default, dti, FUN =
                                                                    median), y = dti, fill = lo
                                                                    geom_boxplot() +
labs(title = "Default Rate by Debt to Income Ratio",
      x = "Default Yes / No",
      y = "Debt to Income Ratio")
default_dti_plot
```



9. Is there a relationship between bank utility score and loan default?

Findings: No, bank utility scores appears to be spread similarly for default “Yes” and default “No”.

```
default_by_bc <- ggplot(data = loan_data, mapping = aes(x = reorder(loan_default, bc_util, FUN =
                                                         median), y = bc_util, fill = loan_default)) +
  geom_violin() +
  geom_jitter(width = 0.07, alpha = 0.5) +
  labs(title = "Default Rate by Bank Utility Score",
       x = "Default Yes / No",
       y = "Bank Utility Score")
default_by_bc
```



Variable Selection

Random Forest Variable Importance

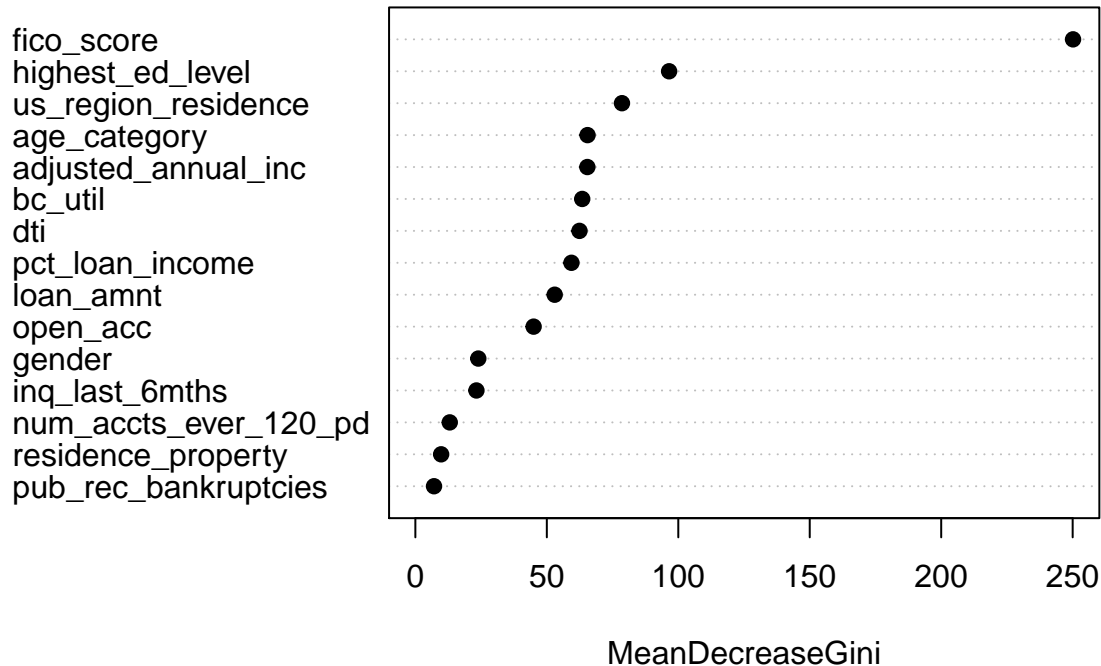
By using the `varImpPlot` function in Random Forests, the most important variables in this data set are determined. The “elbow” method can help eliminate variables that are the least important on the Gini index. After `us_region_residence`, the low variable importance becomes very similar throughout the rest of the list. By eliminating the variables below `us_region_residence`, the model will be simpler and over fitting will be avoided. The variables `fico_score`, `highest_ed_level`, and `us_region_residence` will be used in the predictive models.

```
set.seed(314)

loan_rf <- randomForest(loan_default ~., data = loan_training, importance = TRUE)

varImpPlot(loan_rf, type = 2, pch = 19, main = "Variable Importance in the Loan Data Set")
```

Variable Importance in the Loan Data Set



Predictive Modeling

Random Forests Classification: Predicting loan_default

#First, the model is fit using randomForest() on the training data.
`set.seed(314)`

```
loan_rf_training <- randomForest(loan_default ~ fico_score + highest_ed_level + us_region_residence,
                                data = loan_training, importance = TRUE)
```

#Second, a results table is made.

```
loan_rf_training_results <- data.frame(loan_training,
                                       rforest_pred_0.5 = predict(loan_rf_training,
                                                                newdata = loan_training,
                                                                type = "response"),
                                       predict(loan_rf_training,
                                              newdata = loan_training,
                                              type = "probabilities"))
```

```
loan_rf_training_results %>% dplyr::select(loan_default, rforest_pred_0.5, Yes, No) %>% slice(1:10)
```

##	loan_default	rforest_pred_0.5	Yes	No
## 1	No	No	0.000	1.000
## 2	No	No	0.000	1.000
## 3	No	No	0.010	0.990
## 4	No	No	0.002	0.998
## 5	No	No	0.000	1.000
## 6	Yes	Yes	0.724	0.276
## 7	Yes	Yes	0.894	0.106
## 8	No	No	0.002	0.998
## 9	No	No	0.016	0.984
## 10	No	No	0.000	1.000

The training results table is passed through the confusion matrix function and the optimal cut-off is determined. In this case, the optimal cut is 0.1 with an F1 score of 0.738.

```
cf_matrix(actual_vec = loan_rf_training_results$loan_default,
          pred_prob_vec = loan_rf_training_results$Yes,
          positive_val = "Yes", search_cut = TRUE)
```

##	cut_prob	correct_rate	F1_score	false_pos_rate	false_neg_rate
## 1	0.00	0.2266311	0.3695179	1.0000000000	0.00000000
## 2	0.05	0.8298359	0.7081152	0.1938825851	0.08922559
## 3	0.10	0.8637924	0.7384615	0.1317217563	0.15151515
## 4	0.15	0.8733308	0.7381703	0.1016280217	0.21212121
## 5	0.20	0.8790538	0.7315834	0.0764676862	0.27272727
## 6	0.25	0.8855399	0.7368421	0.0621608288	0.29292929
## 7	0.30	0.8866845	0.7282708	0.0498273310	0.32996633
## 8	0.35	0.8840137	0.7126654	0.0429205723	0.36531987
## 9	0.40	0.8855399	0.7093023	0.0355204736	0.38383838
## 10	0.45	0.8863029	0.7061144	0.0305870745	0.39730640
## 11	0.50	0.8824876	0.6863544	0.0251603355	0.43265993
## 12	0.55	0.8801984	0.6694737	0.0187469166	0.46464646
## 13	0.60	0.8786723	0.6580645	0.0148001973	0.48484848
## 14	0.65	0.8779092	0.6491228	0.0108534780	0.50168350
## 15	0.70	0.8718047	0.6173121	0.0064134188	0.54377104
## 16	0.75	0.8641740	0.5771971	0.0024666996	0.59090909
## 17	0.80	0.8496757	0.5087282	0.0019733596	0.65656566
## 18	0.85	0.8149561	0.3120567	0.0004933399	0.81481481
## 19	0.90	0.7993132	0.2054381	0.0000000000	0.88552189
## 20	0.95	0.7825258	0.0776699	0.0000000000	0.95959596
## 21	1.00	0.7733689	0.0000000	0.0000000000	1.00000000

The random forest model, which was fit on the training data, will now be used on the test data.

#The test results table is made.

```
loan_rf_test_results <- data.frame(loan_test,
                                   rf_pred_0.5 = predict(loan_rf_training,
                                                         newdata = loan_test,
                                                         type = "response"),
                                   predict(loan_rf_training,
                                             newdata = loan_test,
                                             type = "prob"))
loan_rf_test_results %>% dplyr::select(loan_default, rf_pred_0.5, Yes, No) %>% slice(1:10)
```

##	loan_default	rf_pred_0.5	Yes	No
## 1	No	No	0.000	1.000
## 2	No	No	0.044	0.956
## 3	No	No	0.000	1.000
## 4	No	No	0.000	1.000
## 5	No	No	0.034	0.966
## 6	Yes	Yes	0.928	0.072
## 7	No	No	0.000	1.000
## 8	Yes	Yes	0.728	0.272
## 9	No	No	0.066	0.934
## 10	No	No	0.002	0.998

The random forest test results will now be passed through the confusion matrix to determine the F1 score and False Negative observations. The optimal cut-off from the training data set will be used. These numbers

will be compared with the next 2 models.

```
cf_matrix(actual_vec = loan_rf_test_results$loan_default,
           pred_prob_vec = loan_rf_test_results$Yes,
           positive_val = "Yes", cut_prob = .1)
```

```
## $confusion_matrix
##           metric observations      rate pct_total_obs
## 1      Correct           892 0.7935943    0.79359431
## 2 Misclassified           232 0.2064057    0.20640569
## 3 True Positive           152 0.6007905    0.13523132
## 4 True Negative           740 0.8495982    0.65836299
## 5 False Negative          101 0.3992095    0.08985765
## 6 False Positive          131 0.1504018    0.11654804
##
## $F1_summary
##           metric      value
## 1 Precision 0.5371025
## 2 Recall   0.6007905
## 3 F1 Score 0.5671642
```

Results:

Training - F1 = 0.738 False Negative = 0.152

Test - F1 = 0.567 False Negative = 0.399

Decision Tree Classification: Predicting loan_default

#First, the Decision Tree model will be fit on the training data.
`set.seed(314)`

```
loan_tree_training <- rpart(loan_default ~ fico_score + highest_ed_level + us_region_residence,
                           data = loan_training,
                           method = "class",
                           control = rpart.control(cp = 0, minbucket = 4))
```

The results table is created and the optimal cp is found. The optimal cp in the training model is 0.0084.

```
cp_results <- loan_tree_training$cptable %>% data.frame()
round(cp_results, 5)
```

```
##           CP nsplit rel.error  xerror   xstd
## 1  0.24747      0  1.00000 1.00000 0.03608
## 2  0.03451      1  0.75253 0.77778 0.03284
## 3  0.03030      3  0.68350 0.72391 0.03192
## 4  0.01178      5  0.62290 0.67340 0.03099
## 5  0.00842      6  0.61111 0.64310 0.03041
## 6  0.00758      8  0.59428 0.64815 0.03051
## 7  0.00673     10  0.57912 0.64310 0.03041
## 8  0.00505     11  0.57239 0.63636 0.03028
## 9  0.00449     16  0.54545 0.64141 0.03038
## 10 0.00337     19  0.53199 0.64310 0.03041
## 11 0.00224     25  0.51178 0.62795 0.03011
## 12 0.00168     28  0.50505 0.62626 0.03008
## 13 0.00084     42  0.48148 0.64983 0.03054
## 14 0.00056     46  0.47811 0.67003 0.03093
## 15 0.00000     57  0.47138 0.68519 0.03122
```

```
loan_tree_training$cptable
```

```
##          CP nsplit rel error    xerror    xstd
## 1  0.2474747475    0 1.0000000 1.0000000 0.03608279
## 2  0.0345117845    1 0.7525253 0.7777778 0.03284183
## 3  0.0303030303    3 0.6835017 0.7239057 0.03191798
## 4  0.0117845118    5 0.6228956 0.6734007 0.03099448
## 5  0.0084175084    6 0.6111111 0.6430976 0.03041157
## 6  0.0075757576    8 0.5942761 0.6481481 0.03051029
## 7  0.0067340067   10 0.5791246 0.6430976 0.03041157
## 8  0.0050505051   11 0.5723906 0.6363636 0.03027893
## 9  0.0044893378   16 0.5454545 0.6414141 0.03037852
## 10 0.0033670034   19 0.5319865 0.6430976 0.03041157
## 11 0.0022446689   25 0.5117845 0.6279461 0.03011152
## 12 0.0016835017   28 0.5050505 0.6262626 0.03007781
## 13 0.0008417508   42 0.4814815 0.6498316 0.03054305
## 14 0.0005611672   46 0.4781145 0.6700337 0.03093081
## 15 0.0000000000   57 0.4713805 0.6851852 0.03121520
```

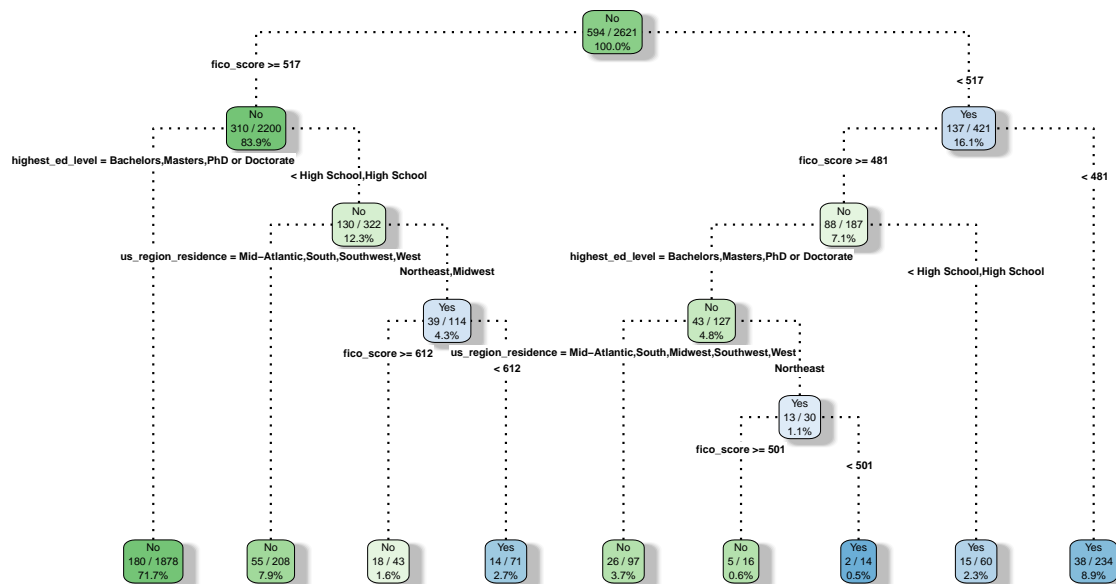
```
cp_results %>% filter(xerror == min(xerror)) %>% mutate(lower_value = xerror - xstd,
  upper_value = xerror + xstd)
```

```
##          CP nsplit rel.error    xerror    xstd lower_value
## 1 0.001683502    28 0.5050505 0.6262626 0.03007781  0.5961848
##   upper_value
## 1  0.6563404
```

#The Decision Tree is pruned using the new cp.

```
loan_pruned <- prune(loan_tree_training, cp = 0.0084)
```

```
rpart.plot(loan_pruned, type = 4, extra = 103, digits = -3,
  box.palette = "GnBu",
  branch.lty = 3, branch.lwd = 3,
  shadow.col = "gray", gap = 0, tweak = 1.0)
```



```
#The Decision Tree results table is made for the training data.
loan_tree_results <- data.frame(loan_training,
                                predict(loan_pruned,
                                         newdata = loan_training,
                                         type = "prob"))

loan_tree_results %>% dplyr::select(loan_default, Yes, No) %>% slice(1:10)
```

	loan_default	Yes	No
## 1	No	0.09584665	0.9041534
## 2	No	0.09584665	0.9041534
## 3	No	0.09584665	0.9041534
## 4	No	0.09584665	0.9041534
## 5	No	0.09584665	0.9041534
## 6	Yes	0.85714286	0.1428571
## 7	Yes	0.75000000	0.2500000
## 8	No	0.09584665	0.9041534
## 9	No	0.09584665	0.9041534
## 10	No	0.09584665	0.9041534

The confusion matrix is used on the results table to determine the F1 score and optimal cut-off. In this case, the F1 score is 0.6457 and the cut-off is 0.375.

```
cf_matrix(actual_vec = loan_tree_results$loan_default,
           pred_prob_vec = loan_tree_results$Yes,
           positive_val = "Yes", search_cut = TRUE)
```

	cut_prob	correct_rate	F1_score	false_pos_rate	false_neg_rate
## 1	0.00	0.2266311	0.36951788	1.0000000000	0.00000000
## 2	0.05	0.2266311	0.36951788	1.0000000000	0.00000000
## 3	0.10	0.8057993	0.61929693	0.1623088308	0.3030303
## 4	0.15	0.8057993	0.61929693	0.1623088308	0.3030303
## 5	0.20	0.8057993	0.61929693	0.1623088308	0.3030303
## 6	0.25	0.8057993	0.61929693	0.1623088308	0.3030303
## 7	0.30	0.8603586	0.64534884	0.0518006907	0.4393939
## 8	0.35	0.8626478	0.64566929	0.0463739517	0.4478114
## 9	0.40	0.8626478	0.64566929	0.0463739517	0.4478114
## 10	0.45	0.8653186	0.63720452	0.0340404539	0.4781145
## 11	0.50	0.8653186	0.63720452	0.0340404539	0.4781145
## 12	0.55	0.8653186	0.63720452	0.0340404539	0.4781145
## 13	0.60	0.8653186	0.63720452	0.0340404539	0.4781145
## 14	0.65	0.8653186	0.63720452	0.0340404539	0.4781145
## 15	0.70	0.8653186	0.63720452	0.0340404539	0.4781145
## 16	0.75	0.8653186	0.63720452	0.0340404539	0.4781145
## 17	0.80	0.8538726	0.58050383	0.0266403552	0.5538721
## 18	0.85	0.7771843	0.03947368	0.0009866798	0.9797980
## 19	0.90	0.7733689	0.00000000	0.0000000000	1.0000000
## 20	0.95	0.7733689	0.00000000	0.0000000000	1.0000000
## 21	1.00	0.7733689	0.00000000	0.0000000000	1.0000000

Now, the Decision Tree model will be fit on the test data.

```
#The results table is created.
loan_tree_test <- data.frame(loan_test,
                              predict(loan_pruned,
                                       newdata = loan_test,
```



```

                                type = "prob"))
loan_tree_test %>% dplyr::select(loan_default, Yes, No) %>% slice(1:10)

##   loan_default      Yes      No
## 1          No 0.09584665 0.9041534
## 2          No 0.09584665 0.9041534
## 3          No 0.09584665 0.9041534
## 4          No 0.09584665 0.9041534
## 5          No 0.09584665 0.9041534
## 6         Yes 0.83760684 0.1623932
## 7          No 0.09584665 0.9041534
## 8         Yes 0.85714286 0.1428571
## 9          No 0.26442308 0.7355769
## 10         No 0.09584665 0.9041534

loan_tree_test <- loan_tree_test %>% mutate(tree_pred_0.3 = ifelse(Yes >= 0.375, "Yes", "No"))
table(loan_tree_test$loan_default, loan_tree_test$tree_pred_0.3)

##
##      No Yes
## No  825  46
## Yes 149 104

```

The confusion matrix will show the F1 score and the count of False Negative observations.

```

cf_matrix(actual_vec = loan_tree_test$loan_default,
           pred_prob_vec = loan_tree_test$Yes,
           positive_val = "Yes", cut_prob = 0.375)

## $confusion_matrix
##      metric observations      rate pct_total_obs
## 1      Correct          929 0.82651246    0.82651246
## 2 Misclassified          195 0.17348754    0.17348754
## 3 True Positive          104 0.41106719    0.09252669
## 4 True Negative          825 0.94718714    0.73398577
## 5 False Negative          149 0.58893281    0.13256228
## 6 False Positive           46 0.05281286    0.04092527
##
## $F1_summary
##      metric      value
## 1 Precision 0.6933333
## 2  Recall 0.4110672
## 3  F1 Score 0.5161290

```

Results:

Training - F1 = 0.646 False Negative = 0.448

Test - F1 = 0.516 False Negative = 0.588

KNN Classification: Predicting loan_default

#First, the optimal k is found. In this case, it is 28.
`set.seed(314)`

```

train.kknn(loan_default ~ fico_score + highest_ed_level + us_region_residence,
            data = loan_training,
            kmax = 40)

```

```
##
## Call:
## train.kknn(formula = loan_default ~ fico_score + highest_ed_level +      us_region_residence, data = 
##
## Type of response variable: nominal
## Minimal misclassification: 0.1320107
## Best kernel: optimal
## Best k: 28
```

#Next, a model is fit on the training data using the optimal k. A results table is also created.
`set.seed(314)`

```
loan_knn_training <- kknn(loan_default ~ fico_score + highest_ed_level + us_region_residence,
                           train = loan_training,
                           test = loan_training,
                           k = 28,
                           distance = 2)
loan_knn_training_results <- data.frame(loan_training,
                                         knn_pred_0.5 = loan_knn_training$fitted.values,
                                         loan_knn_training$prob)

loan_knn_training_results %>% dplyr::select(loan_default, knn_pred_0.5, Yes, No) %>% slice(1:10)
```

	loan_default	knn_pred_0.5	Yes	No
## 1	No	No	0.00000000	1.00000000
## 2	No	No	0.00000000	1.00000000
## 3	No	No	0.00000000	1.00000000
## 4	No	No	0.09869163	0.90130837
## 5	No	No	0.06698121	0.93301879
## 6	Yes	Yes	0.84861473	0.15138527
## 7	Yes	Yes	0.94202921	0.05797079
## 8	No	No	0.11111075	0.88888925
## 9	No	No	0.14406312	0.85593688
## 10	No	No	0.01211420	0.98788580

A confusion matrix is constructed to find the F1 score and optimal cut-off for the KNN model on the training data set. In this case, the F1 score is 0.714 and the optimal cut-off is 0.4.

```
cf_matrix(actual_vec = loan_knn_training_results$loan_default,
           pred_prob_vec = loan_knn_training_results$Yes,
           positive_val = "Yes", search_cut = TRUE)
```

	cut_prob	correct_rate	F1_score	false_pos_rate	false_neg_rate
## 1	0.00	0.2266311	0.36951788	1.00000000	0.00000000
## 2	0.05	0.5757345	0.51652174	0.548593981	0.00000000
## 3	0.10	0.7069821	0.59706191	0.366551554	0.04208754
## 4	0.15	0.7729874	0.64688427	0.269363592	0.08249158
## 5	0.20	0.8119039	0.68131868	0.210162802	0.11279461
## 6	0.25	0.8450973	0.70622287	0.148001973	0.17845118
## 7	0.30	0.8576879	0.70972763	0.115934879	0.23232323
## 8	0.35	0.8698970	0.71223629	0.083374445	0.28956229
## 9	0.40	0.8771461	0.71352313	0.063640849	0.32491582
## 10	0.45	0.8790538	0.70290534	0.048347311	0.36868687
## 11	0.50	0.8748569	0.67779961	0.038973853	0.41919192
## 12	0.55	0.8760015	0.67005076	0.030093735	0.44444444
## 13	0.60	0.8718047	0.64102564	0.020720276	0.49494949

```
## 14      0.65      0.8645555 0.59977452      0.013320178      0.55218855
## 15      0.70      0.8527280 0.54047619      0.009373458      0.61784512
## 16      0.75      0.8405189 0.47355164      0.005920079      0.68350168
## 17      0.80      0.8302175 0.41370224      0.003946719      0.73569024
## 18      0.85      0.8149561 0.31786217      0.001973360      0.80976431
## 19      0.90      0.7977871 0.19452888      0.000000000      0.89225589
## 20      0.95      0.7840519 0.09003215      0.000000000      0.95286195
## 21      1.00      0.7756581 0.02000000      0.000000000      0.98989899
```

Now, the KNN model is fit on the test data set using the optimal k and cut-off from the training set.

```
set.seed(314)

knn_loan_test <- kknnc(loan_default ~ fico_score + highest_ed_level + us_region_residence,
  train = loan_training,
  test = loan_test,
  k = 28, distance = 2)

#The results table is made.
knn_results_test <- data.frame(loan_test,
  knn_pred_0.5 = knn_loan_test$fitted.values,
  knn_loan_test$prob)

knn_results_test <- knn_results_test %>% mutate(knn_pred_0.4 = ifelse(Yes >= 0.4, "Yes", "No"))

knn_results_test %>% dplyr::select(loan_default, knn_pred_0.4, Yes, No) %>% slice(1:10)

##      loan_default knn_pred_0.4      Yes      No
## 1             No             No 0.094580808 0.90541919
## 2             No             No 0.249267718 0.75073228
## 3             No             No 0.000000000 1.00000000
## 4             No             No 0.009844126 0.99015587
## 5             No             No 0.100234086 0.89976591
## 6             Yes             Yes 0.917649374 0.08235063
## 7             No             No 0.000000000 1.00000000
## 8             Yes             Yes 0.846158723 0.15384128
## 9             No             No 0.207715990 0.79228401
## 10            No             No 0.014418174 0.98558183
```

A confusion matrix is made to determine the F1 score and the number of False Negative observations.

```
cf_matrix(actual_vec = knn_results_test$loan_default,
  pred_prob_vec = knn_results_test$Yes,
  positive_val = "Yes",
  cut_prob = 0.4)

## $confusion_matrix
##      metric observations      rate pct_total_obs
## 1      Correct           931 0.82829181      0.82829181
## 2 Misclassified           193 0.17170819      0.17170819
## 3 True Positive           126 0.49802372      0.11209964
## 4 True Negative           805 0.92422503      0.71619217
## 5 False Negative          127 0.50197628      0.11298932
## 6 False Positive           66 0.07577497      0.05871886
##
## $F1_summary
##      metric      value
```

```
## 1 Precision 0.6562500
## 2 Recall 0.4980237
## 3 F1 Score 0.5662921
```

Results:

Training - F1 = 0.714 False Negative = 0.325

Test - F1 = 0.566 False Negative = 0.50

Summary of Findings and Recommendations

Training F1 Score False Negative Rate Random Forest 0.738 0.152 Decision Tree 0.646 0.448 KNN 0.714 0.325

Test F1 Score False Negative Rate False Negative Percent of Observations Random Forest 0.567 0.399 9%
Decision Tree 0.516 0.588 13% KNN 0.566 0.500 11%

The Exploratory Data Analysis was successful in demonstrating what interactions existed in the data set. Findings from the EDA were then supported by the Gini index. Although the EDA showed that variables like `adjusted_annual_inc` and `age_category` have a relationship with `loan_default`, the model appeared to make accurate decisions without these variables. By only using `fico_score`, `highest_ed_level`, and `us_region_residence`, the model remained simple and relatively accurate.

The classification model that provided the best results was the Random Forest.

```
cf_matrix(actual_vec = loan_rf_test_results$loan_default,
           pred_prob_vec = loan_rf_test_results$Yes,
           positive_val = "Yes", cut_prob = .1)
```

```
## $confusion_matrix
##      metric observations      rate pct_total_obs
## 1      Correct           892 0.7935943    0.79359431
## 2 Misclassified           232 0.2064057    0.20640569
## 3 True Positive           152 0.6007905    0.13523132
## 4 True Negative           740 0.8495982    0.65836299
## 5 False Negative          101 0.3992095    0.08985765
## 6 False Positive          131 0.1504018    0.11654804
##
## $F1_summary
##      metric      value
## 1 Precision 0.5371025
## 2 Recall 0.6007905
## 3 F1 Score 0.5671642
```

The Random Forest model has the highest F1 score and the lowest False Negative Rate. As a bank, the situation that would result in the highest risk is predicting that a customer will not default on their loan; however, the customer actually defaults on their loan. This situation can be observed in the model as the False Negative. The Random Forest model has the lowest False Negative Rate. Out of the total observations, the Random Forest model predicted 9% False Negatives. 79% of the observations were classified correctly.

The results of the Random Forest model are a good start for the bank. In order to improve predictions, it would be best if the bank could provide even more data. Perhaps there are also unknown variables that are important in predicting if a customer will default. In the meantime, the bank now has a fairly accurate way to predict default rates.